Eleven scenes, one codec, four λ values; dataset averages against published 3DGS-class compressors

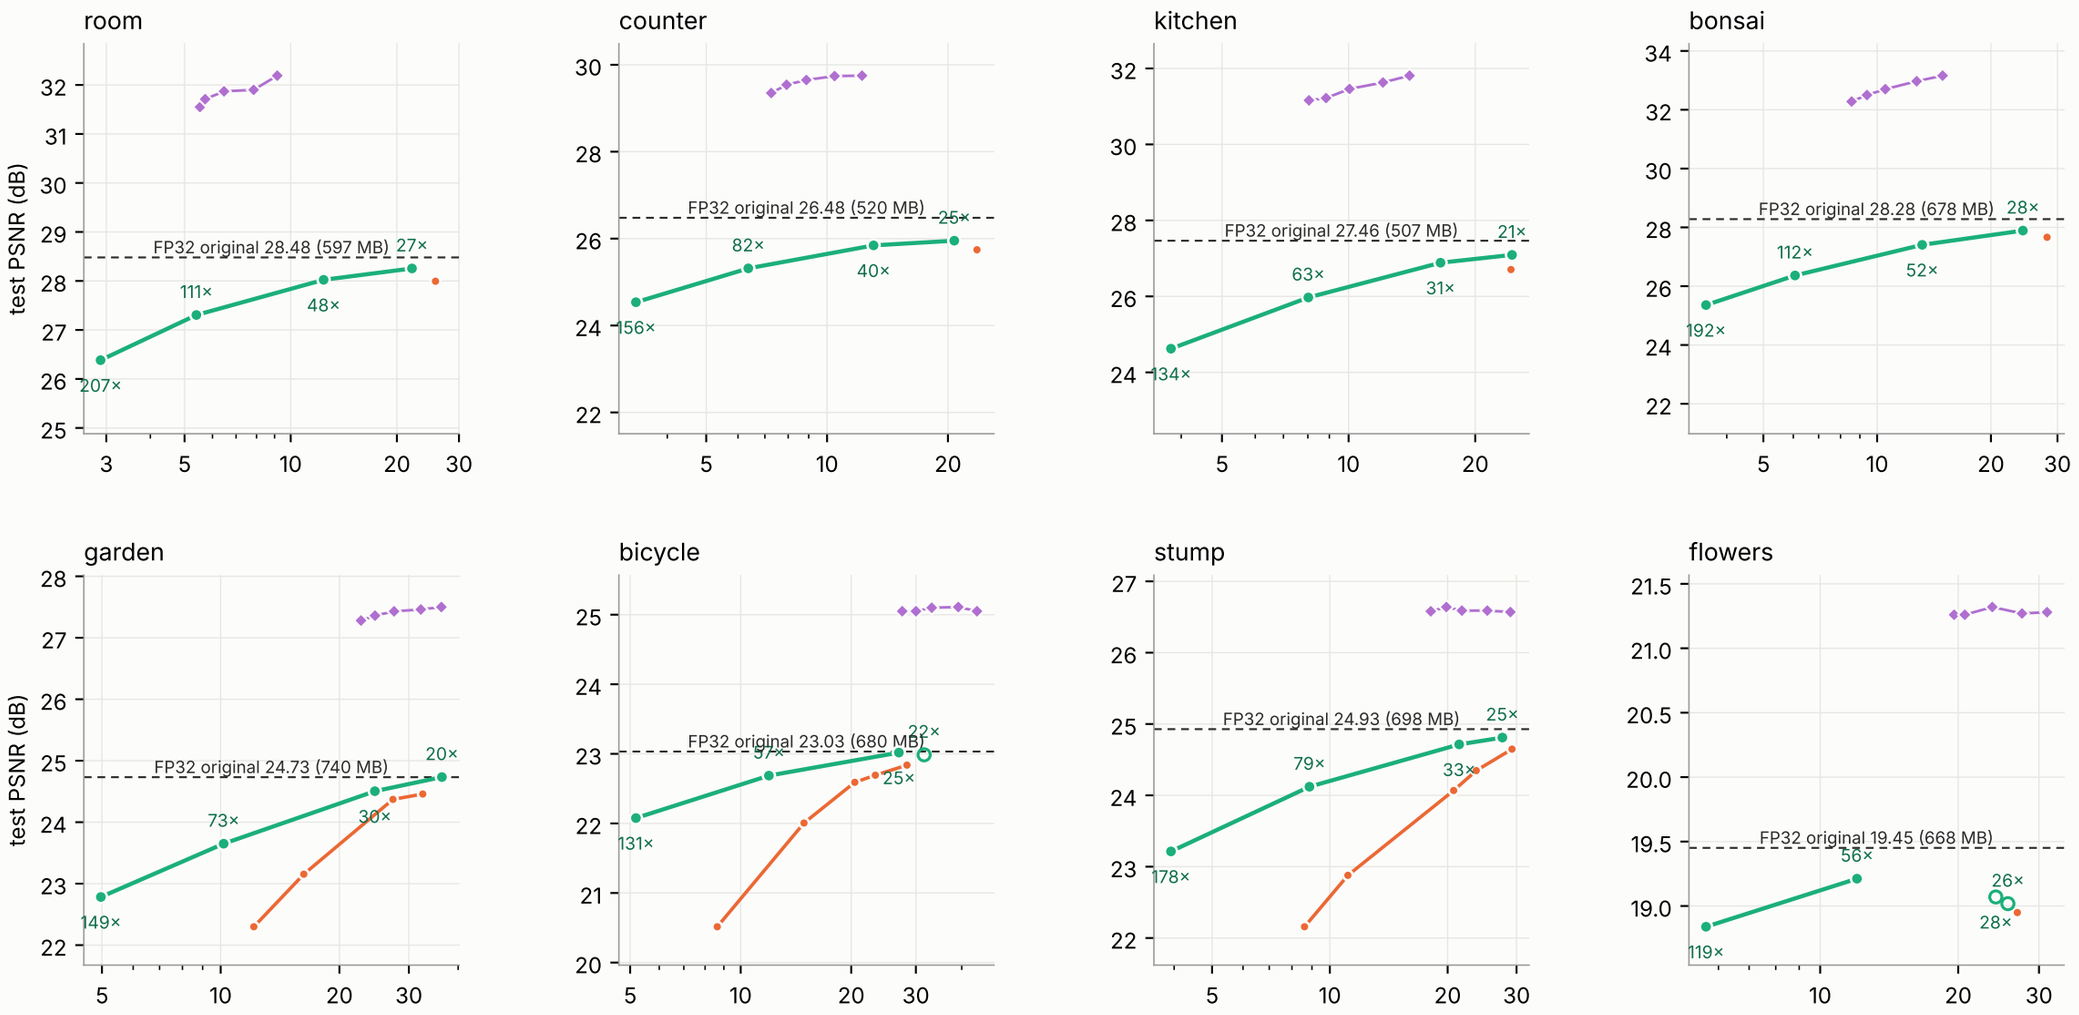

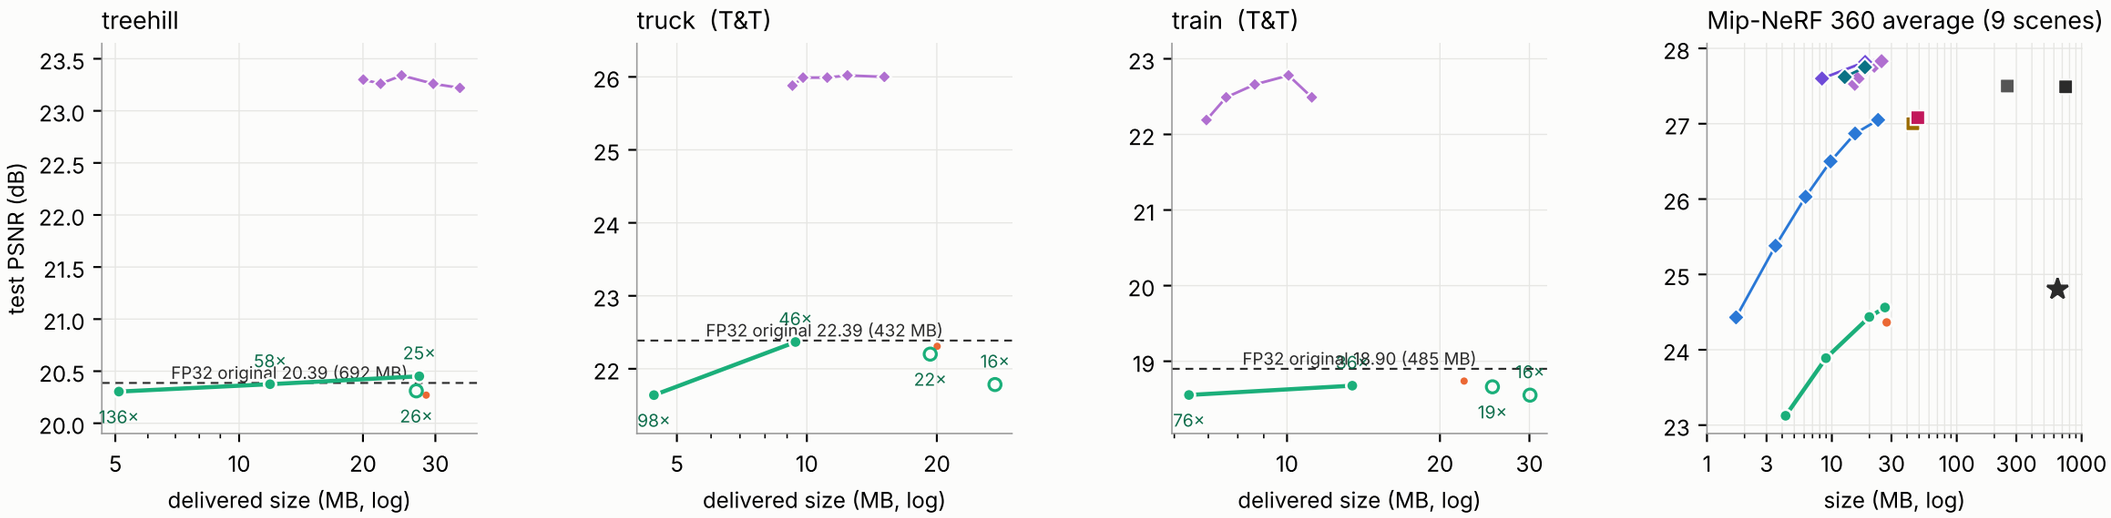

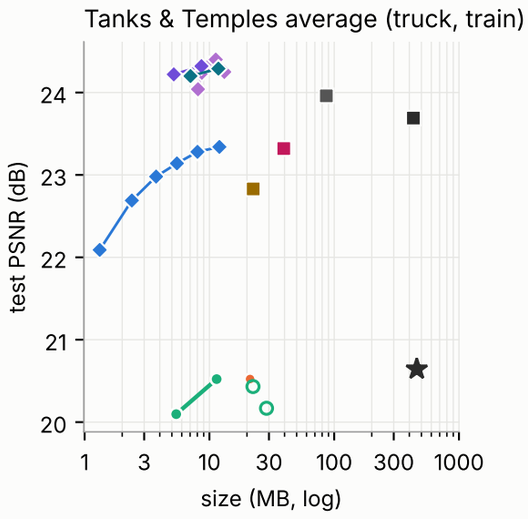

dense training + codec (level 1.0) / + post-hoc collapse
ours: rate-controlled growth + λ-prune (Pareto hull)
HAC
dense training + codec (level 1.0)
FP32 original (642 MB)
3DGS
Scaffold-GS
HAC++
LightGaussian
Compact3DGS
ContextGS
RDO-Gaussian

ours, 11-scene means (MB, ratio vs FP32, Δ test PSNR); indoor scenes = --indoor rerun
λ 6.2e-4: 27.0 MB, 23×, −0.29 dB (Mip-360: −0.24)
λ 1.85e-3: 20.4 MB, 32×, −0.34 dB
λ 6.2e-3: 9.4 MB, 70×, −0.77 dB
λ 1.85e-2: 4.5 MB, 143×, −1.47 dB
level 1.0 (dense + codec): 26.3 MB, 23×, −0.38 dB

3DGS-class numbers are published values: HAC per-scene and average curves (HAC appendix, λe 0.5e-3…4e-3), HAC++ low/high and single points from HAC++ Table II. Per-scene panels show only the HAC curve; the PSNR gap to the 3DGS class is the opaque-mesh representation gap, not the codec.