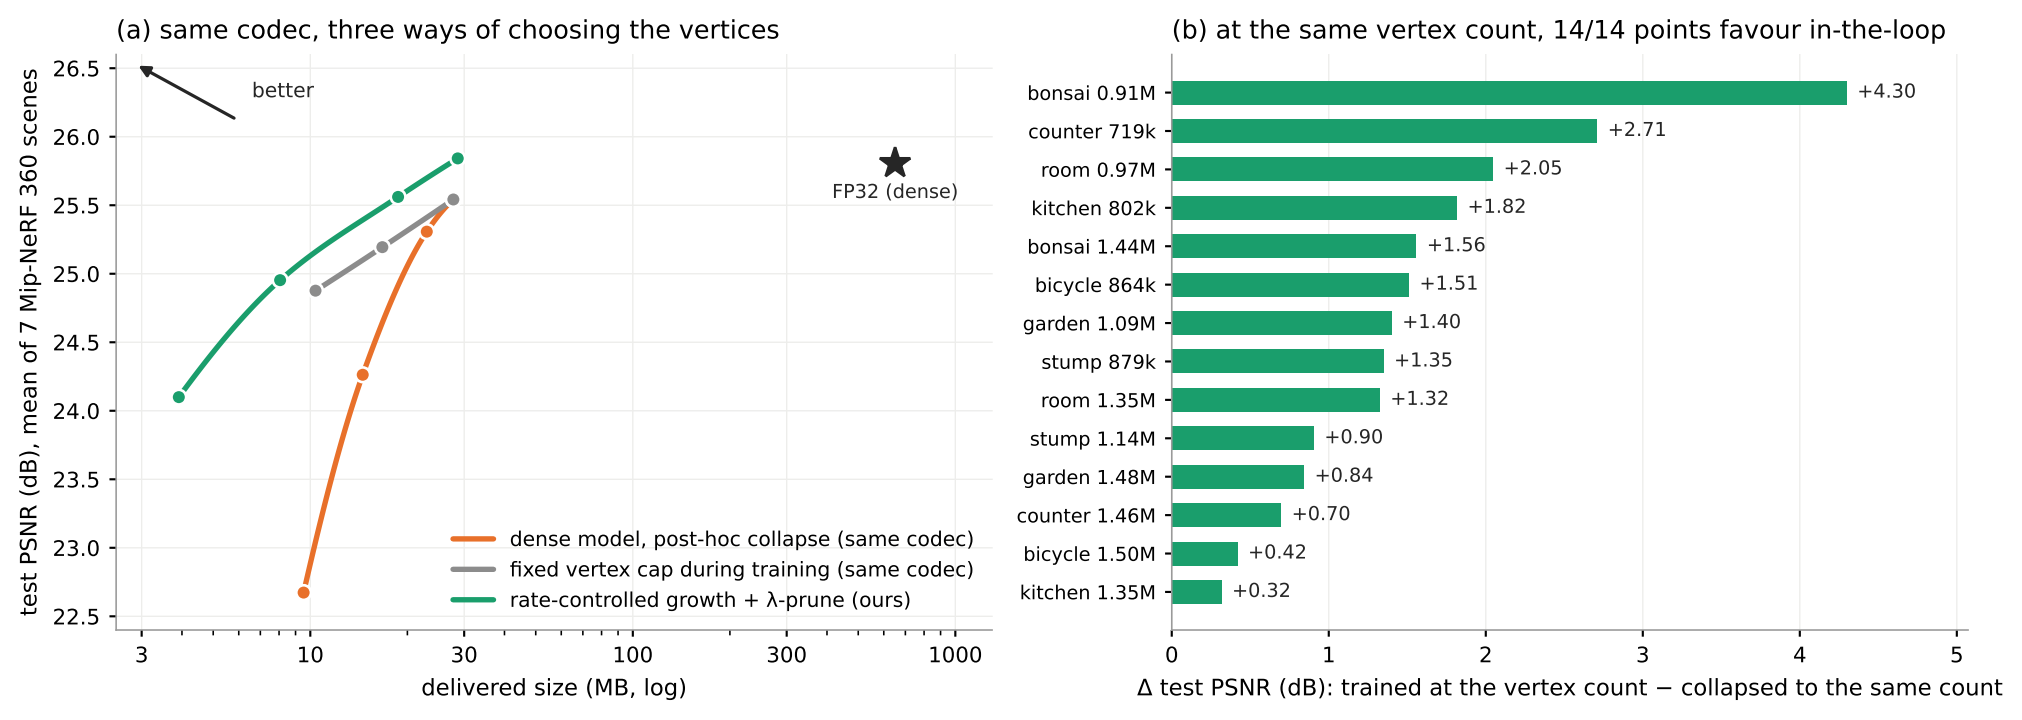
(a) same codec, three ways of choosing the vertices
test PSNR (dB), mean of 7 Mip-NeRF 360 scenes
delivered size (MB, log)
better
FP32 (dense)
dense model, post-hoc collapse (same codec)
fixed vertex cap during training (same codec)
rate-controlled growth + λ-prune (ours)
(b) at the same vertex count, 14/14 points favour in-the-loop
bonsai 0.91M +4.30
counter 719k +2.71
room 0.97M +2.05
kitchen 802k +1.82
bonsai 1.44M +1.56
bicycle 864k +1.51
garden 1.09M +1.40
stump 879k +1.35
room 1.35M +1.32
stump 1.14M +0.90
garden 1.48M +0.84
counter 1.46M +0.70
bicycle 1.50M +0.42
kitchen 1.35M +0.32
Δ test PSNR (dB): trained at the vertex count − collapsed to the same count